


Why Julia?

Hendrik Ranocha

2023-10-09



Scientific computing: simulation of a Mach 3 flow with Trixi.jl¹



Credit: Andrew R. Winters et al. with Trixi.jl


¹R., Schlottke-Lakemper, Winters, Faulhaber, Chan and Gassner (2022); Schlottke-Lakemper, Winters, R. and Gassner (2021)

Numerical analysis: studying time integrators with BSeries.jl²

- ▶ Analysis of numerical integrators for $u'(t) = f(u(t))$

$$u^{n+1} = B(a, \Delta t f, u^n) = a(\emptyset)u^n + \sum_{\tau \in T} \frac{h^{|\tau|}}{\sigma(\tau)} a(\tau) F(\tau)(u^n)$$

- ▶ Based on rooted trees and elementary differentials


$$f(u), f'(f(u)), f'\left(f'(f(u))\right), f''(f(u), f(u)), \dots$$

	Mod. eq.	Mod. int.	Energy pres.
pybs	$\approx 8.3 \text{ s}$	—	$\approx 8.4 \text{ s}$
BSeries.jl	$\approx 0.1 \text{ s}$	$\approx 0.05 \text{ s}$	$\approx 0.1 \text{ s}$

²Ketcheson and R. (2023)

Table of contents

Overview of Julia

Performance

Reproducibility

Caveats

Bringing it all together

Julia

According to the official website

<https://julialang.org>, Julia is ...

- ▶ free
- ▶ dynamic
- ▶ general
- ▶ composable
- ▶ fast
- ▶ reproducible




© Stefan Karpinski / CC BY-NC-SA 4.0

Julia encourages good software development practices!


Testing

- ▶ Testing framework Test.jl in the standard library
- ▶ Continuous integration (CI) via GitHub actions
- ▶ Coverage reports via Coveralls.io and Codecov.io



Documentation

- ▶ Docstrings
- ▶ Doctests and rendering with Documenter.jl
- ▶ GitHub actions and pages



Streamlined release management & dependency tracking via GitHub

JuliaRegistrar.jl Register new versions

```
3      authors = ["Michael Schlottke-Lakemper <mschlott@math.uw.edu>"]
4 - version = "0.3.56-pre"
4 + version = "0.3.56"
5
6 [deps]
7 CodeTracking = "da1fd8a2-8d9e-5ec2-8556-3822fb5608a2"
```

2 comments on commit 39d834e



ranocha replied on Aug 9

@JuliaRegistrar register



JuliaRegistrar replied on Aug 9

Registration pull request created: [JuliaRegistries/General/42462](#)

After the above pull request is merged, it is recommended that a tag is created for this package version.

This will be done automatically if the [Julia TagBot GitHub Action](#) is installed on your GitHub interface, or via:

```
git tag -a v0.3.56 -m "<description of version>" 39d834eb3baai
git push origin v0.3.56
```

TagBot Tag registered versions

v0.3.56

github-actions released this on Aug 9

Trixi v0.3.56

Diff since v0.3.55

Closed issues:

- Include Triangulate.jl as direct dependency? (#754)
- Update StartUpDG.jl to v0.11 (#764)

Merged pull requests:

- Improving performance of DGMulti flux differencing (#757) (@jlchan)
- Astro jet (#772) (@gregorgassner)

Contributors



jlchan and gregorgassner

CompatHelper.jl Keep dependencies up-to-date

CompatHelper: bump compat for "Octavian" to "0.3" #743

Merged

ranocha merged 1 commit into [main](#) from [compathelper/new_version/2021-07-28-00](#) on Jul 28

Conversation 1 Commits 1 Checks 26 Files changed

github-actions bot commented on Jul 28 Contributor

This pull request changes the compat entry for the Octavian package from 0.2.28 to 0.2.28, 0.3.

This keeps the compat entries for earlier versions.

Note: I have not tested your package with this new compat entry. It is your responsibility to make sure that your package tests pass before you merge this pull request.

Trixi.jl developers

Andrés Rueda-Ramírez†, Cologne

methods for plasma simulation

Andrew Winters*, Linköping

shallow water equations

high-order curved meshes

Benedict Geihe†, Cologne

GPU acceleration for HPC

Benjamin Bolm, Cologne

shock capturing methods

Daniel Bach†, Cologne

coupled plasma simulations

Daniel Döhring†, Aachen

time integration methods

David Knapp†, Cologne

adaptive hybrid meshes

application: earth system modeling

Erik Faulhaber, Cologne

particle-based simulation methods

Gregor Gassner*, Cologne

shock-turbulence interaction

Hendrik Ranocha*, Hamburg

adaptive time integration schemes

structure-preserving methods

5 principal developers*, 9 third-party funded scientists†, 4 university-funded scientist

Jesse Chan*, Houston

entropy stable methods

simplex meshes

Johannes Markert†, Cologne

adaptive hybrid meshes

application: tank sloshing

Lars Christmann†, Aachen

high-performance computing

scientific machine learning

Michael Schlottke-Lakemper*, Aachen

adaptive multi-physics simulations

high-performance computing

Niklas Neher†, Stuttgart

particle-based simulation methods

Patrick Ersing, Linköping

shallow water equations

well-balanced schemes

Simon Candelaresi†, Stuttgart

parallel multi-physics coupling

plasma physics

Sophia Schmickler, Cologne

scientific machine learning

Trixi.jl

TrixiParticles.jl

TrixiShallowWater.jl

ReadVTK.jl

Trixi2Vtk.jl

P4est.jl

TrixiBottomTopography.jl

HOHQMesh.jl

SmartShockFinder.jl

Bachelor/master theses

completed: 11

ongoing: 3

PhD theses

completed: 1

ongoing: 7

Interactivity and general programming

Interactivity

- ▶ Julia REPL (read-eval-print-loop)
- ▶ Visual studio code extension
- ▶ Jupyter
- ▶ Pluto.jl notebooks
- ▶ GUIs

General programming

- ▶ Rich type system
- ▶ N-dimensional arrays
- ▶ ...




Documentation: <https://docs.julialang.org>
Type "?" for help, "]?" for Pkg help.
Version 1.9.3 (2023-08-24)
Official <https://julialang.org/> release
julia> █

→ Pluto.jldemo

Example for composability: numerical simulations with uncertainty

Combine


- ▶ Trixi.jl: spatial semidiscretization
- ▶ OrdinaryDiffEq.jl: time integration
- ▶ Measurements.jl: values with uncertainty
velocity = 1.0 ± 0.1
- ▶ Plots.jl: visualization



Credit: R. et al., documentation of Trixi.jl

Serial performance on par with Fortran

- ▶ 3D compressible Euler simulation
(inviscid Taylor-Green vortex)
 - ▶ Curved mesh, entropy-conservative fluxes
 - ▶ Comparable performance as Fortran code FLUXO (same algorithms)
- Performance depends on optimization effort



Parallel scalability experiments on JURECA

JURECA


- ▶ CPU cluster at Jülich Supercomputing Centre, Forschungszentrum Jülich
- ▶ 480 compute nodes
 - ▶ 2xAMD EPYC 7742, 2x64 cores, 2.25 GHz
 - ▶ 512 GB DDR4, 3200 MHz
 - ▶ 128 cores/node, 4 GB/core
- ▶ diskless nodes



Copyright: Forschungszentrum Jülich GmbH / Ralf-Uwe Limbach

Parallel scalability of Trixi.jl


- ▶ Strong scaling experiment on JURECA (MPI only) with Julia v1.8.5
- ▶ 16.8 mio. degrees of freedom
- ▶ Good scalability to 16,384 MPI ranks
- ▶ Comparable performance as Fortran code FLUXO



(higher is better)

Parallel scalability to >50,000 MPI ranks

- ▶ Same setup as before, but 134.2 mio. DOFs
- ▶ Good scalability from 128 to 61,440 cores
- ▶ At full JURECA: 400× speedup,
~34 elements/rank



(higher is better)

Reproducibility


- ▶ Common interfaces allow splitting up tasks and enable code reuse
 - ▶ Depending on others can be scary
 - ▶ Dependency management and reproducibility infrastructure mitigate this
- ▶ Reproducibility is key in modern scientific computing
 - ▶ Dependency management is built into Julia
 - ▶ Binary dependencies are handled as well

“More than 70 % of researchers have tried and failed to reproduce another scientist’s experiments, and more than half have failed to reproduce their own experiments.”

— Baker, Nature 533, 2016

Why should I care about reproducibility in scientific computing?

- ▶ Scientific motivation: best practice/“expected”
- ▶ Legal motivation: my funding agency says so
- ▶ Moral motivation: public money, public X
- ▶ Personal motivation:
Allow others to build upon my results
(and cite me/collaborate with me)



© OpenMoji / CC BY-SA 4.0

Re: Why should I care about reproducibility?

"One of the strengths of this contribution is the accessibility it provides to the algorithms. Computational fluid dynamics packages often involve many underlying dependencies that can take several hours to download, configure, and compile ... By using Julia ..., the authors have significantly reduced this burden: I was able to (begin) reproducing their results within minutes."

— Anonymous Reviewer, ACM TOMS³, 2022


³R., Schlottke-Lakemper, Chan, Rueda-Ramírez, Winters, Hindenlang and Gassner (2021)

Packages

- ▶ Julia packages are Git repositories
- ▶ Package manager `Pkg` in the standard library

```
julia> using Pkg; Pkg.add("Trixi")
```

- ▶ General registry
- ▶ Semantic versioning



Projects

- ▶ Specify direct dependencies with versions
- ▶ Version control friendly version control
(`Project.toml` and `Manifest.toml`)
- ▶ Excellent for reproducible science:
 - ▶ Paper #1: <https://git.io/JYBtP>
 - ▶ Paper #2: <https://git.io/JYBtA>
 - ▶ Paper #3: <https://git.io/JuEIO>
 - ▶ Talk #1: <https://git.io/JqnxE>
 - ▶ Talk #2: <https://git.io/JcLMY>
 - ▶ Talk #3: <https://git.io/JcL6G>
 - ▶ ...

The image shows five GitHub repository cards arranged vertically, each representing a different project:


- README.md**
A purely hyperbolic discontinuous Galerkin approach for self-gravitating gas dynamics
License: MIT DOI: 10.5281/travis.40544769
- README.md**
Preventing pressure oscillations does not fix local linear stability issues of entropy-based split-form high-order schemes
License: MIT DOI: 10.5281/travis.40544769
- README.md**
Julia for adaptive high-order multi-physics simulations
License: MIT Search Issues
- README.md**
Introduction to Julia and Trixi, a numerical simulation framework for hyperbolic PDEs
License: MIT Forked 1,026 Updated 1 day ago Search Issues
- README.md**
JuliaCon 2021: Adaptive and extendable numerical simulations with Trixi.jl
License: MIT Forked 1,026 Updated 1 day ago Search Issues

Binary dependencies

- ▶ Pre-compiled binaries bundled as packages (“JLL” packages)
- ▶ Install via regular package manager
- ▶ Natively call C/Fortran code from Julia

Example: adaptive meshes with p4est

- ▶ Wrapper package P4est.jl
- ▶ Auto-installs binaries with MPI support
- ▶ Works on Linux, macOS, Windows



Calling binaries is fast in Julia⁴

Function Signature	Pybind11	Julia's <code>ccall</code>		Speedup	
int fn0()	132	± 14.9	2.34	± 1.24	56×
int fn1(int)	217	± 20.9	2.35	± 1.33	92×
double fn2(int, double)	232	± 11.7	2.32	± 0.189	100×
char* fn3(int, double, char*)	267	± 28.9	6.27	± 0.396	42×

Table: Round-trip times for calling C functions from Python and Julia in nanoseconds. The benchmark results were collected by using an Intel Core i7-1185G7 CPU running at 3.00 GHz with Julia version 1.7.1, Python version 3.8.10, and Pybind11 version 2.9.1.

⁴Churavy et al. (2022)

BinaryBuilder.jl: create Julia packages with binary artifacts


- ▶ BinaryBuilder.jl: automate building binaries for different targets
- ▶ Cross-compile locally for all Julia-supported hosts
 - ▶ Linux, macOS, Windows, FreeBSD
 - ▶ x86_64, i686, ARM
- ▶ Output: “JLL” package with binary artifacts
- ▶ Yggdrasil: central Julia repo for BB.jl recipes
→ automatically create and register JLLs




Workflow: binary dependency management in Julia


$$\begin{bmatrix} L & A & P & A & C & K \\ L & -A & P & -A & C & -K \\ L & A & P & A & -C & -K \\ L & -A & P & -A & -C & K \\ L & A & -P & -A & C & K \\ L & -A & -P & A & C & -K \end{bmatrix}$$


Workflow: binary dependency management in Julia




Workflow: binary dependency management in Julia



Workflow: binary dependency management in Julia



Workflow: binary dependency management in Julia




Caveat for HPC: using Julia with a system libraries

- ▶ HPC architectures often require using vendor-provided software, e.g., MPI, CUDA
- ▶ Need to replace MPI-enabled JLL binaries by system binaries (e.g., MPI.jl, HDF5.jl)
- ▶ May need to regenerate C bindings for libraries due to MPI ABI change

Remedies:

- ▶ Write/use wrapper packages with corresponding logic
- ▶ Have a look at MPItrampoline
(<https://github.com/eschnett/MPItrampoline>)
- ▶ Document setup procedures for your users and yourself

FAIR data/code principles




© SangyaPundir / CC BY-SA 4.0

- ▶ **Findable:** general registry, JuliaHub, Discourse, Slack, Zulip
- ▶ **Accessible:** open source, package hosting on GitHub, Documenter.jl
- ▶ **Interoperable:** design around informal interfaces, duck typing
- ▶ **Reusable:** package structure and environments, semantic versioning


Memory usage in Julia

- ▶ Julia memory usage has been increasing since Julia v1.6
- ▶ Problematic for MPI-only parallel codes
- ▶ Less problematic for hybrid codes, e.g., MPI+threads or MPI+GPU



Startup latency: challenge for parallel execution

- ▶ Compilation time remains constant
- ▶ Loading time increases for >2000 MPI ranks
- ▶ Julia depot on parallel filesystem (GPFS)
- ▶ Parallel I/O becomes bottleneck




(lower is better)


Julia is promising for scientific computing

Julia hackathon in September

- ▶ C code MAGEMin
(several years of development)
 - ▶ Adaptive meshes with C/C++ code `t8code` wrapped in `T8code.jl`
 - ▶ Parallelization via `Base.Threads`
 - ▶ Coupling, GUI, and visualization in Julia
- 1.5× faster than previous MATLAB version (2 days)



Julia GUI and interface to LaMEM



Credit: Boris Kaus et al.





Open source software Trixi.jl

- ▶ Adaptive high-order simulation framework for conservation laws (MIT license)
- ▶ Goals: usability, extensibility, performance
- ▶ Integration with Julia ecosystem:
 - ▶ OrdinaryDiffEq.jl: time integration
 - ▶ ForwardDiff.jl: automatic differentiation
 - ▶ Plots.jl, Makie.jl: plotting
 - ▶ LoopVectorization.jl: performance
 - ▶ Polyester.jl: multithreading
 - ▶ MPI.jl: distributed parallelism




<https://github.com/trixi-framework/Trixi.jl>

Supersonic flow with curved, adaptive mesh



Credit: Andrew R. Winters et al. with Trixi.jl

Shallow water simulation of Seaside, Oregon, US



Credit: Andrew R. Winters, Sven Goldberg, Maximilian Bertrandt et al. with Trixi.jl

Summary

- ▶ Julia is promising for scientific computing
 - ▶ from laptops to HPC systems
 - ▶ from experimental code to international collaborations
 - ▶ encouraging good research software engineering practices
- ▶ Julia is not perfect
 - ▶ but it's actively developed
 - ▶ it requires some effort to use it well

Summary

- ▶ Julia is promising for scientific computing
 - ▶ from laptops to HPC systems
 - ▶ from experimental code to international collaborations
 - ▶ encouraging good research software engineering practices
- ▶ Julia is not perfect
 - ▶ but it's actively developed
 - ▶ it requires some effort to use it well

<https://ranocha.de>

Thank you

References I

-  Baker, M. (2016). "1,500 scientists lift the lid on reproducibility." In: *Nature* 533, pp. 452–454. doi: 10.1038/533452a.
-  Churavy, V., W. F. Godoy, C. Bauer, H. Ranocha, M. Schlottke-Lakemper, L. Räss, J. Blaschke, M. Giordano, E. Schnetter, S. Omlin, J. S. Vetter, and A. Edelman (Nov. 2022). *Bridging HPC Communities through the Julia Programming Language*. doi: 10.48550/arXiv.2211.02740. arXiv: 2211.02740 [cs.DC].
-  Ketcheson, D. I. and H. Ranocha (June 2023). "Computing with B-series." In: *ACM Transactions on Mathematical Software* 49.2. doi: 10.1145/3573384. arXiv: 2111.11680 [math.NA].
-  Ranocha, H., M. Schlottke-Lakemper, J. Chan, A. M. Rueda-Ramírez, A. R. Winters, F. Hindenlang, and G. J. Gassner (Dec. 2021). *Efficient implementation of modern entropy stable and kinetic energy preserving discontinuous Galerkin methods for conservation laws*. Accepted in ACM Transactions on Mathematical Software (TOMS), 2023. arXiv: 2112.10517 [cs.MS].
-  Ranocha, H., M. Schlottke-Lakemper, A. R. Winters, E. Faulhaber, J. Chan, and G. J. Gassner (Jan. 2022). "Adaptive numerical simulations with Trixi.jl: A case study of Julia for scientific computing." In: *Proceedings of the JuliaCon Conferences* 1.1, p. 77. doi: 10.21105/jcon.00077. arXiv: 2108.06476 [cs.MS].
-  Schlottke-Lakemper, M., A. R. Winters, H. Ranocha, and G. J. Gassner (June 2021). "A purely hyperbolic discontinuous Galerkin approach for self-gravitating gas dynamics." In: *Journal of Computational Physics* 442, p. 110467. doi: 10.1016/j.jcp.2021.110467. arXiv: 2008.10593 [math.NA].